



NASM Tutorial

Yep, it's a tutorial.

CONTENTS

Scope of the Tutorial • Your First Program • Structure of a NASM Program • Details • Your First Few Instructions • The Three Kinds of Operands • Instructions with two memory operands are extremely rare • Defining Data and Reserving Space • Another Example • Using a C Library • Understanding Calling Conventions • Mixing C and Assembly Language • Conditional Instructions • Command Line Arguments • A Longer Example • Floating Point Instructions • Data Sections • Recursion • SIMD Parallelism • Saturated Arithmetic • Graphics • Local Variables and Stack Frames • Using NASM on macOS • Using NASM on Windows • Summary

Scope of the Tutorial

This tutorial will show you how to write assembly language programs on the x86-64 architecture.

You will write both (1) standalone programs and (2) programs that integrate with C.

Don't worry, we won't get too fancy.

Your First Program

Before learning any details, let's make sure you can type in and run programs.

Make sure both `nasm` and `gcc` are installed. Save one of the following programs as **hello.asm**, depending on your machine platform. Then run the program according to the

given instructions.

If you are on a Linux-based OS:

hello.asm

```
; -----
; Writes "Hello, World" to the console using only system calls. Runs on 64-bit Linux
; To assemble and run:
;
;     nasm -felf64 hello.asm && ld hello.o && ./a.out
; -----

        global    _start

_start:  section    .text
        mov     rax, 1          ; system call for write
        mov     rdi, 1          ; file handle 1 is stdout
        mov     rsi, message    ; address of string to output
        mov     rdx, 13         ; number of bytes
        syscall                ; invoke operating system to do the write
        mov     rax, 60         ; system call for exit
        xor     rdi, rdi        ; exit code 0
        syscall                ; invoke operating system to exit

        section    .data
message: db        "Hello, World", 10    ; note the newline at the end
```

```
$ nasm -felf64 hello.asm && ld hello.o && ./a.out
Hello, World
```

If you are on macOS:

hello.asm

```
; -----
; Writes "Hello, World" to the console using only system calls. Runs on 64-bit macOS
; To assemble and run:
;
;     nasm -fmacho64 hello.asm && ld hello.o && ./a.out
; -----

        global    start

start:   section    .text
        mov     rax, 0x02000004    ; system call for write
        mov     rdi, 1          ; file handle 1 is stdout
        mov     rsi, message    ; address of string to output
```

```
mov     rdx, 13           ; number of bytes
syscall                          ; invoke operating system to do the v
mov     rax, 0x02000001    ; system call for exit
xor     rdi, rdi           ; exit code 0
syscall                          ; invoke operating system to exit

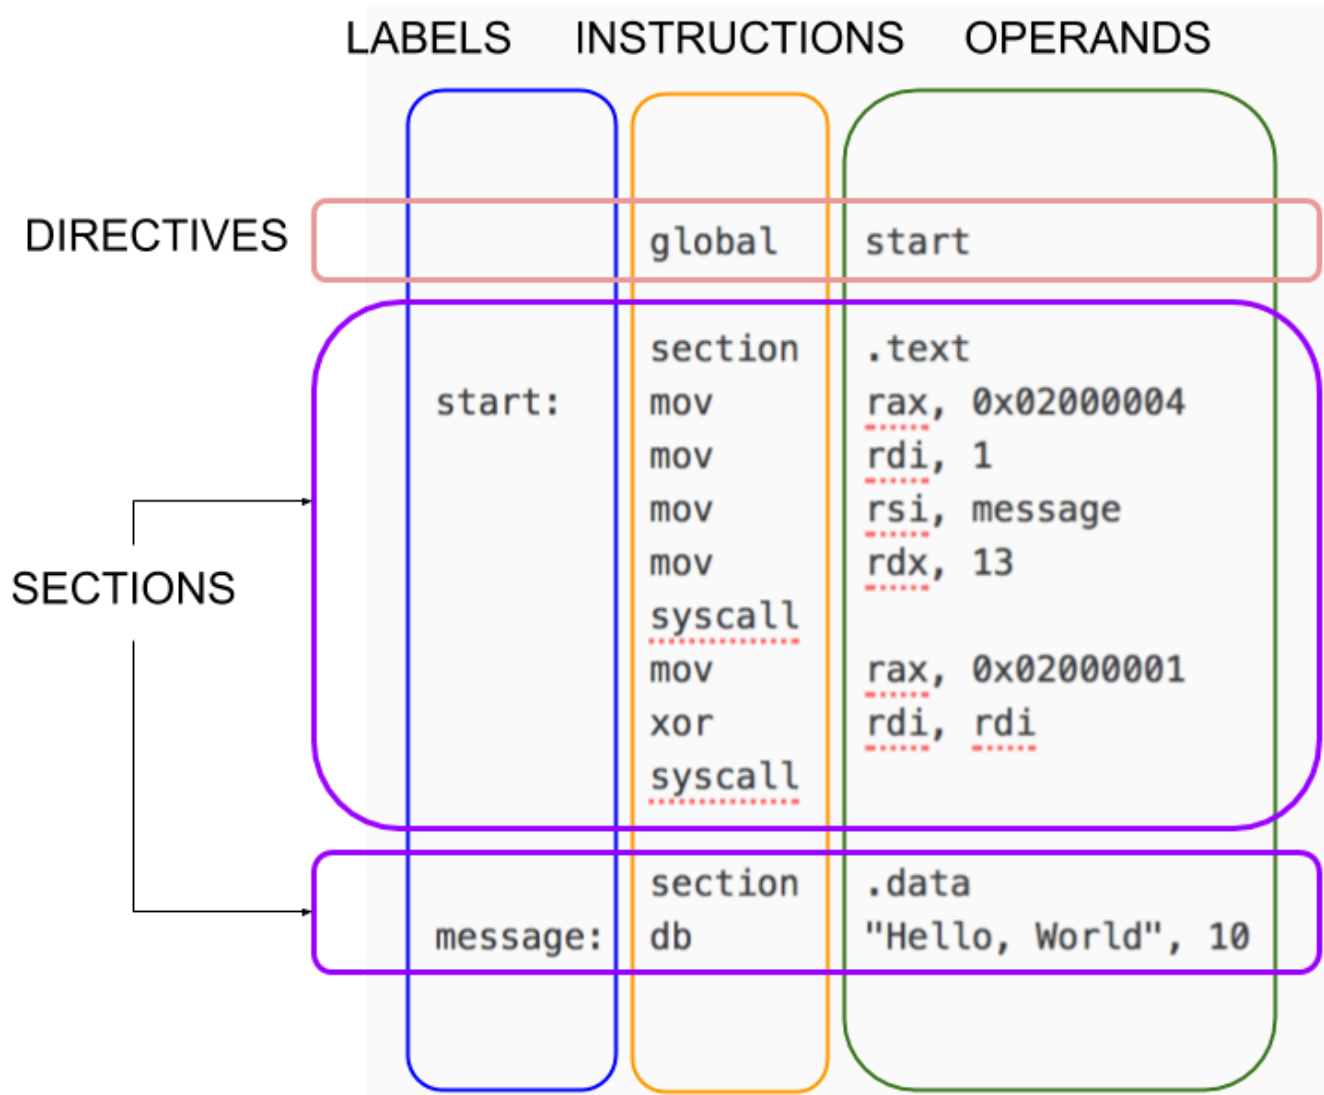
section .data
message: db "Hello, World", 10 ; note the newline at the end
```

```
$ nasm -fmacho64 hello.asm && ld hello.o && ./a.out
Hello, World
```

Exercise: Identify the differences between the two programs.

Structure of a NASM Program

NASM is line-based. Most programs consist of **directives** followed by one or more **sections**. Lines can have an optional **label**. Most lines have an **instruction** followed by zero or more **operands**.



Generally, you put code in a section called `.text` and your constant data in a section called `.data`.

Details

NASM is an awesome assembler, but assembly language is complex. You need more than a tutorial. You need details. Lots of details. Be ready to consult:

- [The NASM Manual](#), which is pretty good!
- [The Intel Processor Manuals](#)

Your First Few Instructions

There are hundreds of instructions. You can't learn them all at once. Just start with these:

| Instruction | Description |
|--------------------|-------------------------------------------------------------------------------------------------------|
| mov x, y | $x \leftarrow y$ |
| and x, y | $x \leftarrow x \wedge y$ |
| or x, y | $x \leftarrow x \vee y$ |
| xor x, y | $x \leftarrow x \oplus y$ |
| add x, y | $x \leftarrow x + y$ |
| sub x, y | $x \leftarrow x - y$ |
| inc x | $x \leftarrow x + 1$ |
| dec x | $x \leftarrow x - 1$ |
| syscall n | Invoke operating system routine n |
| db | A pseudo-instruction that declares bytes that will be in memory when the program runs |

The Three Kinds of Operands

It helps to know these cold. Get ready to memorize them. How can you remember things? [Nicky Case can help you with this!](#)

Register Operands

In this tutorial we only care about the integer registers, the flag register, and the xmm registers. (If you are familiar with the x86 architecture, you will know that this means we are skipping the FP, MMX, YMM, segment, control, debug, test, and protected mode registers.) Hopefully you have already been introduced to the x86-64 architecture, in which case, this is a quick review. The 16 integer registers are 64 bits wide and are called:

| | | | | | | | | | | | | | | | |
|------------------|------------------|------------------|------------------|------------------|------------------|------------------|------------------|----|----|-----|-----|-----|-----|-----|-----|
| R0 aka RAX | R1 aka RCX | R2 aka RDX | R3 aka RBX | R4 aka RSP | R5 aka RBP | R6 aka RSI | R7 aka RDI | R8 | R9 | R10 | R11 | R12 | R13 | R14 | R15 |
|------------------|------------------|------------------|------------------|------------------|------------------|------------------|------------------|----|----|-----|-----|-----|-----|-----|-----|

(Note that 8 of the registers have alternate names.) You can treat the lowest 32-bits of each register as a register itself but using these names:

| | | | | | | | | | | | | | | |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|------|------|------|------|------|
| R0D | R1D | R2D | R3D | R4D | R5D | R6D | R7D | R8D | R9D | R10D | R11D | R12D | R13D | R14D |
| aka | aka | aka | aka | aka | aka | aka | aka | | | | | | | |
| EAX | ECX | EDX | EBX | ESP | EBP | ESI | EDI | | | | | | | |
| R15D | | | | | | | | | | | | | | |

You can treat the lowest 16-bits of each register as a register itself but using these names:

| | | | | | | | | | | | | | | |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|------|------|------|------|------|
| R0W | R1W | R2W | R3W | R4W | R5W | R6W | R7W | R8W | R9W | R10W | R11W | R12W | R13W | R14W |
| aka | aka | aka | aka | aka | aka | aka | aka | | | | | | | |
| AX | CX | DX | BX | SP | BP | SI | DI | | | | | | | |
| R15W | | | | | | | | | | | | | | |

You can treat the lowest 8-bits of each register as a register itself but using these names:

| | | | | | | | | | | | | | | |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|------|------|------|------|------|
| R0B | R1B | R2B | R3B | R4B | R5B | R6B | R7B | R8B | R9B | R10B | R11B | R12B | R13B | R14B |
| aka | aka | aka | aka | aka | aka | aka | aka | | | | | | | |
| AL | CL | DL | BL | SPL | BPL | SIL | DIL | | | | | | | |
| R15B | | | | | | | | | | | | | | |

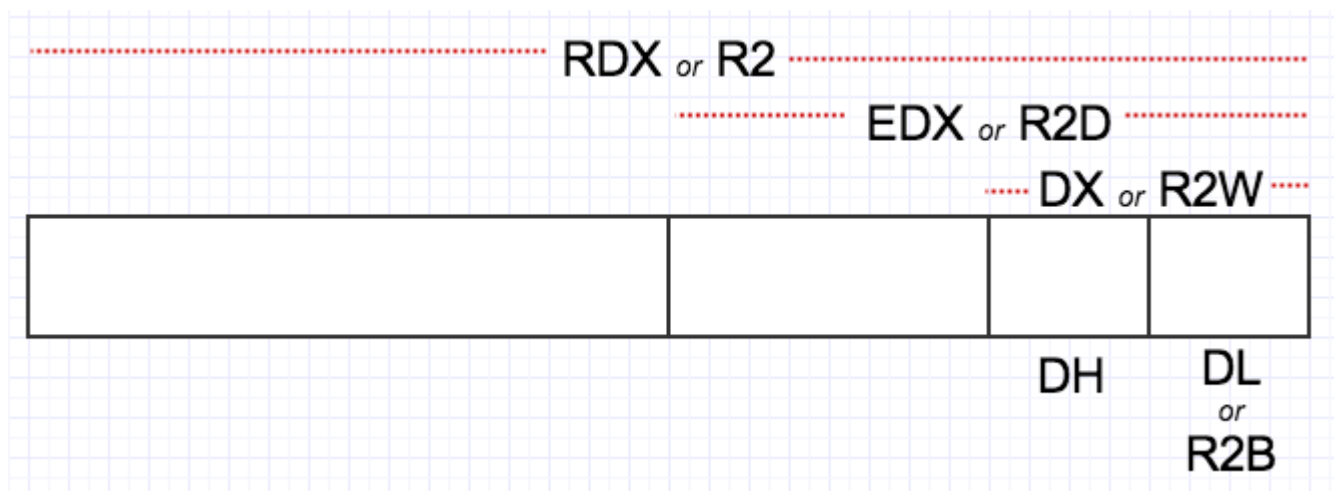
For historical reasons, bits 15 through 8 of R0..R3 are named:

| | | | |
|----|----|----|----|
| AH | CH | DH | BH |
|----|----|----|----|

And finally, there are 16 XMM registers, each 128 bits wide, named:

| | | | | | | | | | | | |
|-------|-------|-------|-------|------|------|------|------|------|------|-------|-------|
| XMM0 | XMM1 | XMM2 | XMM3 | XMM4 | XMM5 | XMM6 | XMM7 | XMM8 | XMM9 | XMM10 | XMM11 |
| XMM12 | XMM13 | XMM14 | XMM15 | | | | | | | | |

Study this picture; hopefully it helps:



Memory Operands

These are the basic forms of addressing:

- [number]
- [reg]
- [reg + reg*scale] scale is 1, 2, 4, or 8 only
- [reg + number]
- [reg + reg*scale + number]

The number is called the **displacement**; the plain register is called the **base**; the register with the scale is called the **index**.

Examples:

```
[750]           ; displacement only
[rbp]           ; base register only
[rcx + rsi*4]    ; base + index * scale
[rbp + rdx]      ; scale is 1
[rbx - 8]        ; displacement is -8
[rax + rdi*8 + 500] ; all four components
[rbx + counter]  ; uses the address of the variable 'counter' as the displacement
```

Immediate Operands

These can be written in many ways. Here are some examples from the official docs.

```
200           ; decimal
0200          ; still decimal - the leading 0 does not make it octal
0200d         ; explicitly decimal - d suffix
```

```
0d200      ; also decimal - 0d prefix
0c8h       ; hex - h suffix, but leading 0 is required because c8h looks like a
0xc8       ; hex - the classic 0x prefix
0hc8       ; hex - for some reason NASM likes 0h
310q       ; octal - q suffix
0q310      ; octal - 0q prefix
11001000b  ; binary - b suffix
0b1100_1000 ; binary - 0b prefix, and by the way, underscores are allowed
```

Instructions with two memory operands are extremely rare

In fact, we'll not see any such instruction in this tutorial. Most of the basic instructions have only the following forms:

- `add reg, reg`
- `add reg, mem`
- `add reg, imm`
- `add mem, reg`
- `add mem, imm`

Defining Data and Reserving Space

These examples come from [Chapter 3 of the docs](#). To place data in memory:

```
db 0x55      ; just the byte 0x55
db 0x55,0x56,0x57 ; three bytes in succession
db 'a',0x55   ; character constants are OK
db 'hello',13,10,'$' ; so are string constants
dw 0x1234     ; 0x34 0x12
dw 'a'        ; 0x61 0x00 (it's just a number)
dw 'ab'       ; 0x61 0x62 (character constant)
dw 'abc'      ; 0x61 0x62 0x63 0x00 (string)
dd 0x12345678 ; 0x78 0x56 0x34 0x12
dd 1.234567e20 ; floating-point constant
```



```

dq    0x123456789abcdef0    ; eight byte constant
dq    1.234567e20            ; double-precision float
dt    1.234567e20            ; extended-precision float

```

There are other forms; check the NASM docs. Later.

To reserve space (without initializing), you can use the following pseudo instructions. They should go in a section called `.bss` (you'll get an error if you try to use them in a `.text` section):

```

buffer:      resb    64            ; reserve 64 bytes
wordvar:     resw    1            ; reserve a word
realarray:   resq    10           ; array of ten reals

```

Another Example

Here's a macOS program to study:

triangle.asm

```

; -----
; This is an macOS console program that writes a little triangle of asterisks to
; output. Runs on macOS only.
;
;    nasm -fmacho64 triangle.asm && gcc hola.o && ./a.out
; -----

        global  start
        section .text

start:

        mov     rdx, output            ; rdx holds address of next byte to v
        mov     r8, 1                  ; initial line length
        mov     r9, 0                  ; number of stars written on line so

line:

        mov     byte [rdx], '*'        ; write single star
        inc     rdx                    ; advance pointer to next cell to wri
        inc     r9                     ; "count" number so far on line
        cmp     r9, r8                 ; did we reach the number of stars fo
        jne     line                  ; not yet, keep writing on this line

lineDone:

        mov     byte [rdx], 10         ; write a new line char
        inc     rdx                    ; and move pointer to where next char
        inc     r8                     ; next line will be one char longer

```

```

        mov     r9, 0                ; reset count of stars written on this line
        cmp     r8, maxlines         ; wait, did we already finish the last line?
        jng     line                 ; if not, begin writing this line

done:
        mov     rax, 0x02000004       ; system call for write
        mov     rdi, 1                ; file handle 1 is stdout
        mov     rsi, output           ; address of string to output
        mov     rdx, dataSize         ; number of bytes
        syscall                       ; invoke operating system to do the write
        mov     rax, 0x02000001       ; system call for exit
        xor     rdi, rdi              ; exit code 0
        syscall                       ; invoke operating system to exit

        section .bss
maxlines equ 8
dataSize equ 44
output:  resb dataSize

```

```

$ nasm -fmacho64 triangle.asm && ld triangle.o && ./a.out
*
**
***
****
*****
*****
*****
*****
*****

```

New things in this example:

- `cmp` does a comparison
- `je` jumps to a label if the previous comparison was equal. We also have `jne` (jump if not equal), `jl` (jump if less), `jnl` (jump if not less), `jg` (jump if greater), `jng` (jump if not greater), `jle` (jump if less or equal), `jnl` (jump if not less or equal), `jge` (jump if greater or equal), `jnge` (jump if not greater or equal), and many more.
- `equ` is actually not a real instruction. It simply defines an abbreviation for the assembler itself to use. (This is a profound idea.)
- The `.bss` section is for *writable* data.

Using a C Library

Writing standalone programs with just system calls is cool, but rare. We would like to use the good stuff in the C library.

Remember how in C execution “starts” at the function `main`? That’s because the C library actually has the `_start` label inside itself! The code at `_start` does some initialization, then it calls `main`, then it does some clean up, then it issues the system call for exit. So you just have to implement `main`. We can do that in assembly!

If you have Linux, try this:

hola.asm

```
; -----
; Writes "Hola, mundo" to the console using a C library. Runs on Linux.
;
;   nasm -felf64 hola.asm && gcc hola.o && ./a.out
; -----

        global    main
        extern    puts

        section   .text

main:                                ; This is called by the C library start
        mov      rdi, message        ; First integer (or pointer) argument
        call     puts                ; puts(message)
        ret                               ; Return from main back into C library

message:
        db       "Hola, mundo", 0    ; Note strings must be terminated with \0
```

```
$ nasm -felf64 hola.asm && gcc hola.o && ./a.out
Hola, mundo
```

Under macOS, it will look a little different:

hola.asm

```
; -----
; This is an macOS console program that writes "Hola, mundo" on one line and then
; It uses puts from the C library. To assemble and run:
;
;   nasm -fmacho64 hola.asm && gcc hola.o && ./a.out
; -----

        global    _main
        extern    _puts

        section   .text

_main:
        push     rbx                  ; Call stack must be aligned
        lea      rdi, [rel message]  ; First argument is address of message
        call     _puts                ; puts(message)
        pop      rbx                  ; Fix up stack before returning
```

```

ret

section .data
message: db "Hola, mundo", 0 ; C strings need a zero byte at the e

```

```

$ nasm -f macho64 hola.asm && gcc hola.o && ./a.out
Hola, mundo

```

In macOS land, C functions (or any function that is exported from one module to another, really) must be prefixed with underscores. The call stack must be aligned on a 16-byte boundary (more on this later). And when accessing named variables, a `rel` prefix is required.

Understanding Calling Conventions

How did we know the argument to `puts` was supposed to go in `RDI`? Answer: there are a number of conventions that are followed regarding calls.

When writing code for 64-bit Linux that integrates with a C library, you must follow the calling conventions explained in the [AMD64 ABI Reference](#). You can also get this information from [Wikipedia](#). The most important points are:

- From left to right, pass as many parameters as will fit in registers. The order in which registers are allocated, are:
 - For integers and pointers, `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`.
 - For floating-point (float, double), `xmm0`, `xmm1`, `xmm2`, `xmm3`, `xmm4`, `xmm5`, `xmm6`, `xmm7`.
- Additional parameters are pushed on the stack, right to left, and are to be *removed by the caller* after the call.
- After the parameters are pushed, the call instruction is made, so when the called function gets control, the return address is at `[rsp]`, the first memory parameter is at `[rsp+8]`, etc.
- **The stack pointer `rsp` must be aligned to a 16-byte boundary before making a call.** Fine, but the process of making a call pushes the return address (8 bytes) on

the stack, so when a function gets control, `rsp` is not aligned. You have to make that extra space yourself, by pushing something or subtracting 8 from `rsp`.

- The only registers that the called function is required to preserve (the callee-save registers) are: `rbp`, `rbx`, `r12`, `r13`, `r14`, `r15`. All others are free to be changed by the called function.
- The callee is also supposed to save the control bits of the XMMCSR and the x87 control word, but x87 instructions are rare in 64-bit code so you probably don't have to worry about this.
- Integers are returned in `rax` or `rdx:rax`, and floating point values are returned in `xmm0` or `xmm1:xmm0`.

Got that? No? What's need is more examples, and practice.

Here is a program that illustrates how registers have to be saved and restored:

fib.asm

```

; -----
; A 64-bit Linux application that writes the first 90 Fibonacci numbers. To
; assemble and run:
;
;     nasm -felf64 fib.asm && gcc fib.o && ./a.out
; -----

    global  main
    extern  printf

    section .text

main:
    push    rbx                ; we have to save this since we use it

    mov     ecx, 90             ; ecx will countdown to 0
    xor     rax, rax            ; rax will hold the current number
    xor     rbx, rbx            ; rbx will hold the next number
    inc     rbx                ; rbx is originally 1

print:
    ; We need to call printf, but we are using rax, rbx, and rcx.  printf
    ; may destroy rax and rcx so we will save these before the call and
    ; restore them afterwards.

    push    rax                ; caller-save register
    push    rcx                ; caller-save register

    mov     rdi, format         ; set 1st parameter (format)
    mov     rsi, rax            ; set 2nd parameter (current_number)
    xor     rax, rax            ; because printf is varargs

```

```

; Stack is already aligned because we pushed three 8 byte registers
call    printf                ; printf(format, current_number)

pop     rcx                    ; restore caller-save register
pop     rax                    ; restore caller-save register

mov     rdx, rax               ; save the current number
mov     rax, rbx               ; next number is now current
add     rbx, rdx               ; get the new next number
dec     ecx                    ; count down
jnz     print                  ; if not done counting, do some more

pop     rbx                    ; restore rbx before returning
ret

format:
db     "%20ld", 10, 0

```

```

$ nasm -felf64 fib.asm && gcc fib.o && ./a.out
0
1
1
2
.
.
.
679891637638612258
1100087778366101931
1779979416004714189

```

We just saw some new instructions:

| Instruction | Description |
|--------------------------|----------------------------------------------------------------------------------------------------|
| push <i>x</i> | Decrement <code>rsp</code> by the size of the operand, then store <i>x</i> in <code>[rsp]</code> |
| pop <i>x</i> | Move <code>[rsp]</code> into <i>x</i> , then increment <code>rsp</code> by the size of the operand |
| jnz <i>label</i> | If the processor's Z (zero) flag is set, jump to the given label |
| call <i>label</i> | Push the address of the next instruction, then jump to the label |
| ret | Pop into the instruction pointer |

Mixing C and Assembly Language

This program is just a simple function that takes in three integer parameters and returns the maximum value.

maxofthree.asm

```
; -----
; A 64-bit function that returns the maximum value of its three 64-bit integer
; arguments. The function has signature:
;
; int64_t maxofthree(int64_t x, int64_t y, int64_t z)
;
; Note that the parameters have already been passed in rdi, rsi, and rdx. We
; just have to return the value in rax.
; -----

global maxofthree
section .text
maxofthree:
    mov     rax, rdi                ; result (rax) initially holds x
    cmp     rax, rsi                ; is x less than y?
    cmovl   rax, rsi                ; if so, set result to y
    cmp     rax, rdx                ; is max(x,y) less than z?
    cmovl   rax, rdx                ; if so, set result to z
    ret                               ; the max will be in rax
```

Here is a C program that calls the assembly language function.

callmaxofthree.c

```
/*
 * A small program that illustrates how to call the maxofthree function we wrote
 * assembly language.
 */

#include <stdio.h>
#include <inttypes.h>

int64_t maxofthree(int64_t, int64_t, int64_t);

int main() {
    printf("%ld\n", maxofthree(1, -4, -7));
    printf("%ld\n", maxofthree(2, -6, 1));
    printf("%ld\n", maxofthree(2, 3, 1));
    printf("%ld\n", maxofthree(-2, 4, 3));
    printf("%ld\n", maxofthree(2, -6, 5));
    printf("%ld\n", maxofthree(2, 4, 6));
    return 0;
}
```

```
$ nasm -felf64 maxofthree.asm && gcc callmaxofthree.c maxofthree.o && ./a.out
1
2
3
4
5
6
```

Conditional Instructions

After an arithmetic or logic instruction, or the compare instruction, `cmp`, the processor sets or clears bits in its `rflags`. The most interesting flags are:

- `s` (sign)
- `z` (zero)
- `c` (carry)
- `o` (overflow)

So after doing, say, an addition instruction, we can perform a jump, move, or set, based on the new flag settings. For example:

| Instruction | Description |
|----------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>jz <i>L</i></code> | Jump to label <i>L</i> if the result of the operation was zero |
| <code>cmovno <i>x</i>, <i>y</i></code> | $x \leftarrow y$ if the last operation did not overflow |
| <code>setc <i>x</i></code> | $x \leftarrow 1$ if the last operation had a carry, but $x \leftarrow 0$ otherwise (<i>x</i> must be a byte-size register or memory location) |

The conditional instructions have three base forms: `j` for conditional jump, `cmov` for conditional move, and `set` for conditional set. The suffix of the instruction has one of the 30 forms: `s ns z nz c nc o no p np pe po e ne l nl le nle g ng ge nge a na ae nae b nb be nbe`.

Command Line Arguments

You know that in C, `main` is just a plain old function, and it has a couple parameters of its own:

```
int main(int argc, char** argv)
```

So, you guessed it, `argc` will end up in `rdi`, and `argv` (a pointer) will end up in `rsi`. Here is a program that uses this fact to simply echo the command line arguments to a program, one per line:

echo.asm

```
; -----
; A 64-bit program that displays its command line arguments, one per line.
;
; On entry, rdi will contain argc and rsi will contain argv.
; -----

        global  main
        extern  puts
        section .text

main:
        push    rdi                ; save registers that puts uses
        push    rsi
        sub     rsp, 8             ; must align stack before call

        mov     rdi, [rsi]         ; the argument string to display
        call    puts              ; print it

        add     rsp, 8             ; restore %rsp to pre-aligned value
        pop     rsi               ; restore registers puts used
        pop     rdi

        add     rsi, 8             ; point to next argument
        dec     rdi               ; count down
        jnz     main              ; if not done counting keep going

        ret
```

```
$ nasm -felf64 echo.asm && gcc echo.o && ./a.out dog 22 -zzz "hi there"
./a.out
dog
22
-zzz
hi there
```

A Longer Example

Note that as far as the C Library is concerned, command line arguments are always strings. If you want to treat them as integers, call `atoi`. Here's a neat program to compute x^y .

power.asm

```
; -----
; A 64-bit command line application to compute x^y.
;
; Syntax: power x y
; x and y are (32-bit) integers
; -----

global main
extern printf
extern puts
extern atoi

section .text
main:
    push    r12                ; save callee-save registers
    push    r13
    push    r14
    ; By pushing 3 registers our stack is already aligned for calls

    cmp     rdi, 3              ; must have exactly two arguments
    jne     error1

    mov     r12, rsi            ; argv

    ; We will use ecx to count down from the exponent to zero, esi to hold the
    ; value of the base, and eax to hold the running product.

    mov     rdi, [r12+16]      ; argv[2]
    call    atoi                ; y in eax
    cmp     eax, 0              ; disallow negative exponents
    jl      error2
    mov     r13d, eax           ; y in r13d

    mov     rdi, [r12+8]       ; argv
    call    atoi                ; x in eax
    mov     r14d, eax           ; x in r14d

    mov     eax, 1              ; start with answer = 1

check:
```

```
test    r13d, r13d           ; we're counting y down to 0
jz      gotit                ; done
imul    eax, r14d            ; multiply in another x
dec     r13d
jmp     check

gotit:                                ; print report on success
mov     rdi, answer
movsxd  rsi, eax
xor     rax, rax
call    printf
jmp     done

error1:                                ; print error message
mov     edi, badArgumentCount
call    puts
jmp     done

error2:                                ; print error message
mov     edi, negativeExponent
call    puts

done:                                ; restore saved registers
pop     r14
pop     r13
pop     r12
ret

answer:
db      "%d", 10, 0
badArgumentCount:
db      "Requires exactly two arguments", 10, 0
negativeExponent:
db      "The exponent may not be negative", 10, 0
```

```
$ nasm -felf64 power.asm && gcc -o power power.o
$ ./power 2 19
524288
$ ./power 3 -8
The exponent may not be negative
$ ./power 1 500
1
$ ./power 1
Requires exactly two arguments
```

Floating Point Instructions

Floating-point arguments go into the xmm registers. Here is a simple function for summing the values in a double array:

sum.asm

```

; -----
; A 64-bit function that returns the sum of the elements in a floating-point
; array. The function has prototype:
;
; double sum(double[] array, uint64_t length)
; -----

        global sum
        section .text

sum:
        xorpd    xmm0, xmm0           ; initialize the sum to 0
        cmp      rsi, 0                ; special case for length = 0
        je       done
next:
        addsd    xmm0, [rdi]          ; add in the current array element
        add      rdi, 8                ; move to next array element
        dec      rsi                  ; count down
        jnz      next                ; if not done counting, continue
done:
        ret                          ; return value already in xmm0

```

Note the floating point instructions have an `sd` suffix; that's the most common one, but we'll see some other ones later. Here is a C program that calls it:

callsum.c

```

/*
 * Illustrates how to call the sum function we wrote in assembly language.
 */

#include <stdio.h>
#include <inttypes.h>

double sum(double[], uint64_t);

int main() {
    double test[] = {
        40.5, 26.7, 21.9, 1.5, -40.5, -23.4
    };
    printf("%20.7f\n", sum(test, 6));
    printf("%20.7f\n", sum(test, 2));
    printf("%20.7f\n", sum(test, 0));
    printf("%20.7f\n", sum(test, 3));
    return 0;
}

```

```
$ nasm -felf64 sum.asm && gcc sum.o callsum.c && ./a.out
26.70000000
67.20000000
0.00000000
89.10000000
```

Data Sections

The text section is read-only on most operating systems, so you might find the need for a data section. On most operating systems, the data section is only for initialized data, and you have a special .bss section for uninitialized data. Here is a program that averages the command line arguments, expected to be integers, and displays the result as a floating point number.

average.asm

```
; -----
; 64-bit program that treats all its command line arguments as integers and
; displays their average as a floating point number. This program uses a data
; section to store intermediate results, not that it has to, but only to
; illustrate how data sections are used.
; -----

global    main
extern    atoi
extern    printf
default   rel

section   .text
main:
    dec    rdi                ; argc-1, since we don't count program name
    jz     nothingToAverage
    mov    [count], rdi       ; save number of real arguments
accumulate:
    push   rdi                ; save register across call to atoi
    push   rsi
    mov    rdi, [rsi+rdi*8]    ; argv[rdi]
    call   atoi                ; now rax has the int value of arg
    pop    rsi                 ; restore registers after atoi call
    pop    rdi
    add    [sum], rax          ; accumulate sum as we go
    dec    rdi                 ; count down
    jnz    accumulate         ; more arguments?
average:
```

```
    cvtsi2sd xmm0, [sum]
    cvtsi2sd xmm1, [count]
    divsd     xmm0, xmm1           ; xmm0 is sum/count
    mov      rdi, format          ; 1st arg to printf
    mov      rax, 1               ; printf is varargs, there is 1 non-int arg

    sub      rsp, 8               ; align stack pointer
    call     printf               ; printf(format, sum/count)
    add      rsp, 8               ; restore stack pointer

    ret

nothingToAverage:
    mov      rdi, error
    xor      rax, rax
    call     printf
    ret

section .data
count: dq 0
sum:    dq 0
format: db "%g", 10, 0
error:  db "There are no command line arguments to average", 10, 0
```

```
$ nasm -felf64 average.asm && gcc average.o && ./a.out 19 8 21 -33
3.75
$ nasm -felf64 average.asm && gcc average.o && ./a.out
There are no command line arguments to average
```

This program highlighted some processor instructions that convert between integers and floating point values. A few of the most common are:

| Instruction | Description |
|------------------------------------------|-------------------------------------------------------|
| cvtsi2sd <i>xmmreg, r / m32</i> | <i>xmmreg</i> [63..0] ← intToDouble(<i>r / m32</i>) |
| cvtsi2ss <i>xmmreg, r / m32</i> | <i>xmmreg</i> [31..0] ← intToFloat(<i>r / m32</i>) |
| cvtsd2si <i>reg32, xmmr / m64</i> | <i>reg32</i> ← doubleToInt(<i>xmmr / m64</i>) |
| cvtss2si <i>reg32, xmmr / m32</i> | <i>reg32</i> ← floatToInt(<i>xmmr / m32</i>) |

Recursion

Perhaps surprisingly, there's nothing out of the ordinary required to implement recursive functions. You just have to be careful to save registers, as usual. Pushing and popping around the recursive call is a typical strategy.

factorial.asm

```
; -----
; An implementation of the recursive function:
;
;  uint64_t factorial(uint64_t n) {
;      return (n <= 1) ? 1 : n * factorial(n-1);
;  }
; -----

        global  factorial

        section .text
factorial:
        cmp     rdi, 1           ; n <= 1?
        jnbe    L1              ; if not, go do a recursive call
        mov     rax, 1           ; otherwise return 1
        ret

L1:
        push    rdi              ; save n on stack (also aligns %rsp!)
        dec     rdi              ; n-1
        call    factorial        ; factorial(n-1), result goes in %rax
        pop     rdi              ; restore n
        imul    rax, rdi         ; n * factorial(n-1), stored in %rax
        ret
```

An example caller:

callfactorial.c

```
/*
 * An application that illustrates calling the factorial function defined elsewhere
 */

#include <stdio.h>
#include <inttypes.h>

uint64_t factorial(uint64_t n);

int main() {
    for (uint64_t i = 0; i < 20; i++) {
        printf("factorial(%2lu) = %lu\n", i, factorial(i));
    }
    return 0;
}
```

```
$ nasm -felf64 factorial.asm && gcc -std=c99 factorial.o callfactorial.c && ./a.out
factorial( 0) = 1
factorial( 1) = 1
factorial( 2) = 2
factorial( 3) = 6
factorial( 4) = 24
factorial( 5) = 120
factorial( 6) = 720
factorial( 7) = 5040
factorial( 8) = 40320
factorial( 9) = 362880
factorial(10) = 3628800
factorial(11) = 39916800
factorial(12) = 479001600
factorial(13) = 6227020800
factorial(14) = 87178291200
factorial(15) = 1307674368000
factorial(16) = 20922789888000
factorial(17) = 355687428096000
factorial(18) = 6402373705728000
factorial(19) = 121645100408832000
```

SIMD Parallelism

The XMM registers can do arithmetic on floating point values one operation at a time (scalar) or multiple operations at a time (packed). The operations have the form:

```
op xmmreg_or_memory, xmmreg
```

For floating point addition, the instructions are:

| Instruction | Description |
|--------------|----------------------------------------------------------------------------------------------------------|
| addpd | Do two double-precision additions in parallel (add packed double) |
| addsd | Do just one double-precision addition, using the low 64-bits of the register (add scalar double) |
| addps | Do four single-precision additions in parallel (add packed single) |
| addss | Do just one single-precision addition, using the low 32-bits of the register (add scalar single) |

Here’s a function that adds four floats at once:

```
add_four_floats.asm
```

```
; void add_four_floats(float x[4], float y[4])
; x[i] += y[i] for i in range(0..4)
```



```
global add_four_floats
section .text

add_four_floats:
    movdqa xmm0, [rdi]           ; all four values of x
    movdqa xmm1, [rsi]           ; all four values of y
    addps  xmm0, xmm1            ; do all four sums in one shot
    movdqa [rdi], xmm0
    ret
```

and a caller:

```
#include <stdio.h>
void add_four_floats(float[], float[]);

int main() {
    float x[] = {-29.750, 244.333, 887.29, 48.1E22};
    float y[] = {29.750, 199.333, -8.29, 22.1E23};
    add_four_floats(x, y);
    printf("%f\n%f\n%f\n%f\n", x[0], x[1], x[2], x[3]);
    return 0;
}
```

Also see this [nice little x86 floating-point slide deck from Ray Seyfarth](#).

Saturated Arithmetic

The XMM registers can also do arithmetic on integers. The instructions have the form:

op xmmreg_or_memory, xmmreg

For integer addition, the instructions are:

| Instruction | Description |
|--------------|----------------------|
| paddb | Do 16 byte-additions |
| paddw | Do 8 word-additions |
| paddd | Do 4 dword-additions |
| paddq | Do 2 qword-additions |

| | |
|----------------|---------------------------------------------------------|
| paddsb | Do 16 byte-additions with signed saturation (80..7F) |
| paddsw | Do 8 word-additions with signed saturation (8000..7F) |
| paddusb | Do 16 byte-additions with unsigned saturation (00..FF) |
| paddusw | Do 8 word-additions with unsigned saturation (00..FFFF) |

Here's an example. It also illustrates how you load the XMM registers. You can't load immediate values; you have to use `movaps` to move from memory. There are other ways, but we're not covering everything in this tutorial.

satexample.asm

```
; -----
; Example of signed saturated arithmetic.
; -----

        global  main
        extern  printf

        section .text
main:
        push    rbp
        movaps  xmm0, [arg1]
        movaps  xmm1, [arg2]
        paddsw  xmm0, xmm1
        movaps  [result], xmm0

        lea     rdi, [format]
        mov     esi, dword [result]
        mov     edx, dword [result+4]
        mov     ecx, dword [result+8]
        mov     r8d, dword [result+12]
        xor     rax, rax
        call    printf
        pop     rbp
        ret

        section .data
        align   16
arg1:    dw      0x3544, 0x24FF, 0x7654, 0x9A77, 0xF677, 0x9000, 0xFFFF, 0x0000
arg2:    dw      0x7000, 0x1000, 0xC000, 0x1000, 0xB000, 0xA000, 0x1000, 0x0000
result:  dd      0, 0, 0, 0
format:  db      '%X%X%X%X', 10, 0
```

Graphics

Any C program can be “ported” to assembly language. That goes for graphics programs, too.



This program probably does not work.

I last tested this in 2003. Back in the old-school OpenGL days. Used Win32. Pre-GLSL days. Used GLUT. I haven't had access to a Windows box in a while and I'm not even sure it will work anymore. This is presented here for historical interest only. If you can modify it to work under modern OpenGL, please let me know. I'll update the program and cite your contribution, of course!

triangle.asm

```
; -----
; triangle.asm
;
; A very simple *Windows* OpenGL application using the GLUT library. It
; draws a nicely colored triangle in a top-level application window. One
; interesting thing is that the Windows GL and GLUT functions do NOT use the
; C calling convention; instead they use the "stdcall" convention which is
; like C except that the callee pops the parameters.
; -----

global _main
extern _glClear@4
extern _glBegin@4
extern _glEnd@0
extern _glColor3f@12
extern _glVertex3f@12
extern _glFlush@0
extern _glutInit@8
extern _glutInitDisplayMode@4
extern _glutInitWindowPosition@8
extern _glutInitWindowSize@8
extern _glutCreateWindow@4
extern _glutDisplayFunc@4
extern _glutMainLoop@0

section .text
title: db 'A Simple Triangle', 0
zero: dd 0.0
one: dd 1.0
half: dd 0.5
neghalf: dd -0.5

display:
push dword 16384
```

```

call    _glClear@4           ; glClear(GL_COLOR_BUFFER_BIT)
push    dword 9
call    _glBegin@4           ; glBegin(GL_POLYGON)
push    dword 0
push    dword 0
push    dword [one]
call    _glColor3f@12        ; glColor3f(1, 0, 0)
push    dword 0
push    dword [neghalf]
push    dword [neghalf]
call    _glVertex3f@12       ; glVertex(-.5, -.5, 0)
push    dword 0
push    dword [one]
push    dword 0
call    _glColor3f@12        ; glColor3f(0, 1, 0)
push    dword 0
push    dword [neghalf]
push    dword [half]
call    _glVertex3f@12       ; glVertex(.5, -.5, 0)
push    dword [one]
push    dword 0
push    dword 0
call    _glColor3f@12        ; glColor3f(0, 0, 1)
push    dword 0
push    dword [half]
push    dword 0
call    _glVertex3f@12       ; glVertex(0, .5, 0)
call    _glEnd@0             ; glEnd()
call    _glFlush@0           ; glFlush()
ret

```

_main:

```

push    dword [esp+8]        ; push argv
lea     eax, [esp+8]         ; get addr of argc (offset changed :-)
push    eax
call    _glutInit@8          ; glutInit(&argc, argv)
push    dword 0
call    _glutInitDisplayMode@4
push    dword 80
push    dword 80
call    _glutInitWindowPosition@8
push    dword 300
push    dword 400
call    _glutInitWindowSize@8
push    title
call    _glutCreateWindow@4
push    display
call    _glutDisplayFunc@4
call    _glutMainLoop@0
ret

```

Local Variables and Stack Frames


First, please read [Eli Bendersky's article](#). That overview is more complete than my brief notes.

When a function is called the caller will first put the parameters in the correct registers then issue the `call` instruction. Additional parameters beyond those covered by the registers will be pushed on the stack prior to the call. The call instruction puts the return address on the top of stack. So if you have the function:

```
int64_t example(int64_t x, int64_t y) {
    int64_t a, b, c;
    b = 7;
    return x * b + y;
}
```

Then on entry to the function, *x* will be in edi, *y* will be in esi, and the return address will be on the top of the stack. Where can we put the local variables? An easy choice is on the stack itself, though if you have enough registers, use them! Registers tend to be faster anyway.

If you are running on a machine that respect the standard ABI, you can leave `rsp` where it is and access the “extra parameters” and the local variables directly from `rsp` for example:



So our function looks like this:

```

global example
section .text
example:
    mov     qword [rsp-16], 7
    mov     rax, rdi
    imul    rax, [rsp+8]
    add     rax, rsi
    ret

```

If our function were to make another call, you would have to adjust `rsp` to get out of the way at that time.

On Windows you can't use this scheme because if an interrupt were to occur, everything above the stack pointer gets plastered. This doesn't happen on most other operating systems because there is a **red zone** of 128 bytes past the stack pointer which is safe from these things. In this case, you can make room on the stack immediately:

```

example:
    sub     rsp, 24

```

so our stack looks like this:

| | |
|--------|----------------------------|
| rsp | a |
| rsp+8 | b |
| rsp+16 | c |
| rsp+24 | retaddr |
| rsp+32 | caller's stack frame |

Here's the function now. Note that we have to remember to replace the stack pointer before returning!

```

global example
section .text
example:
    sub     rsp, 24
    mov     qword [rsp+8], 7
    mov     rax, rdi
    imul    rax, [rsp+8]

```

```

add    rax, rsi
add    rsp, 24
ret

```

Using NASM on macOS

Hopefully you've gone through the whole tutorial above using a Linux-based operating system (or perhaps more correctly, and ELF64 system). There are pretty much only five things to know to get these examples working under a 64-bit macOS system:

- This object file format is `macho64`, not `elf64`.
- The system call numbers are *totally different*.
- Symbols shared between modules will be prefixed by underscores.
- It seems that the gcc linker in macOS doesn't allow absolute addressing unless you tweak some settings. So add `default rel` when you are referencing labeled memory locations, and always use `lea` to get your addresses.
- Also, it appears that sometimes under Linux, the 16-bit stack alignment requirement is not enforced, but it appears to be *always* enforced under macOS.

So here's the average program from above, written for macOS.

average.asm

```

; -----
; 64-bit program that treats all its command line arguments as integers and
; displays their average as a floating point number. This program uses a data
; section to store intermediate results, not that it has to, but only to
; illustrate how data sections are used.
;
; Designed for OS X. To assemble and run:
;
;     nasm -fmacho64 average.asm && gcc average.o && ./a.out
; -----

global  _main
extern  _atoi
extern  _printf
default rel

section .text
_main:
    push    rbx                ; we don't ever use this, but it is neces
                                ; to align the stack so we can call stuff

```

```

    dec     rdi                ; argc-1, since we don't count program name
    jz      nothingToAverage
    mov     [count], rdi      ; save number of real arguments

accumulate:
    push    rdi                ; save register across call to atoi
    push    rsi                ; save register across call to atoi
    mov     rdi, [rsi+rdi*8]    ; argv[rdi]
    call    _atoi             ; now rax has the int value of arg
    pop     rsi                ; restore registers after atoi call
    pop     rdi
    add     [sum], rax         ; accumulate sum as we go
    dec     rdi                ; count down
    jnz     accumulate        ; more arguments?

average:
    cvtsi2sd xmm0, [sum]
    cvtsi2sd xmm1, [count]
    divsd   xmm0, xmm1        ; xmm0 is sum/count
    lea     rdi, [format]      ; 1st arg to printf
    mov     rax, 1             ; printf is varargs, there is 1 non-int arg
    call    _printf            ; printf(format, sum/count)
    jmp     done

nothingToAverage:
    lea     rdi, [error]
    xor     rax, rax
    call    _printf

done:
    pop     rbx                ; undoes the stupid push at the beginning
    ret

section .data
count: dq 0
sum: dq 0
format: db "%g", 10, 0
error: db "There are no command line arguments to average", 10, 0

```

```

$ nasm -fmacho64 average.asm && gcc average.o && ./a.out
There are no command line arguments to average
$ nasm -fmacho64 average.asm && gcc average.o && ./a.out 54.3
54
$ nasm -fmacho64 average.asm && gcc average.o && ./a.out 54.3 -4 -3 -25 455.1111
95.4

```

Using NASM on Windows

I'm not sure what the system calls are on Windows, but I do know that if you want to assemble and link with the C library, you have to understand [the x64 conventions](#). Read them. You will learn such things as:

- The first four integer parameters are passed in RCX, RDX, R8, and R9. The rest are to be pushed on the stack.
- The callee must preserve RBX, RBP, RDI, RSI, RSP, R12, R13, R14, and R15.
- The first four floating point parameters are passed in, you guessed it, XMM0, XMM1, XMM2, and XMM3.
- Return values go in RAX or XMM0.

IMPORTANT: There's one thing that's really hard to find in any documentation: the x64 calling convention requires you to allocate 32 bytes of [shadow space](#) before each call, and remove it after your call. This means your "hello world" program looks like this:

hello.asm

```
; -----
; This is a Win64 console program that writes "Hello" on one line and then exits.
; uses puts from the C library. To assemble and run:
;
;     nasm -fwin64 hello.asm && gcc hello.obj && a
; -----

global  main
extern  puts
section .text

main:
    sub     rsp, 28h                ; Reserve the shadow space
    mov     rcx, message           ; First argument is address of message
    call    puts                   ; puts(message)
    add     rsp, 28h                ; Remove shadow space
    ret

message:
    db      'Hello', 0             ; C strings need a zero byte at the end
```

Did you notice we actually reserved 40 bytes? Thirty-two bytes of shadow space is a minimum requirement. In our `main` function, we are calling another function, so our stack [must be aligned on a 16-byte boundary](#). When `main` is called, the return address (8 bytes) was pushed, so we have to "add" an extra 8 bytes to the shadow space.

Summary

We've covered:

- ✓ How to run a NASM program
- ✓ The structure of a NASM program
- ✓ The most basic instructions
- ✓ Instruction formats
- ✓ Mixing C and assembly language
- ✓ Floating point, saturated arithmetic, and parallel instructions
- ✓ Calls, calling conventions, and recursion
- ✓ Some platform-specific details